

Graduate School Class Reminders

- ▶ Maintain six feet of distancing
- ▶ Please sit in the same chair each class time
- ▶ Observe entry/exit doors as marked
- ▶ Use hand sanitizer when you enter/exit the classroom
- ▶ Use a disinfectant wipe/spray to wipe down your learning space before and after class
- ▶ Media Services: 414 955-4357 option 2

Documentation on the web

- ▶ CRAN: <http://cran.r-project.org>
- ▶ R manuals: <https://cran.r-project.org/manuals.html>
- ▶ SAS: <http://support.sas.com/documentation>
- ▶ Step-by-Step Programming with Base SAS 9.4 (SbS):
<https://documentation.sas.com/api/docsets/basess/9.4/content/basess.pdf>
- ▶ SAS 9.4 Programmer's Guide: Essentials (PGE):
<https://documentation.sas.com/api/docsets/lepg/9.4/content/lepg.pdf>
- ▶ Wiki: <https://wiki.biostat.mcw.edu> (MCW/VPN)

HW 2: ISO 8601 and the Proleptic Gregorian Calendar

- ▶ ISO 8601 dictates that the Proleptic Gregorian Calendar be used for dates prior to the start of the Gregorian Calendar
- ▶ Is the R Date class ISO 8601 compliant?
- ▶ If not, then at what date does it diverge?
- ▶ Unix/R define 1970-01-01 as date 0 counting for-/back-wards
- ▶ The Julian Period covers all of written human history starting on Monday 4713BCE-01-01 which is ISO 8601 date -2,440,588
- ▶ What Year-Month-Day is this in the corresponding Proleptic Gregorian Calendar?
- ▶ Hints: do this without for loops
the R interpreter will take a long time for 2.4M iterations
using `data.frames` and vectorization this is pretty fast
for example, in `julian.R` see my implementation of the Proleptic Julian Calendar
- ▶ Use two `data.frames`: one for each Calendar and combine them (but you can't use `merge` for 2.4M records)
- ▶ Use `sprintf` to concatenate character vectors

Calendars, dates and ISO 8601: see Wikipedia

- ▶ The Date class represents dates
- ▶ The Egyptian Calendar had 365 days/year circa 3000BCE
- ▶ Leap days were added haphazardly as seen fit
- ▶ Emperor Julius Caesar solidified the calendar
- ▶ The Julian Calendar created a leap day once every 4 years (with some initial issues and corresponding corrections)
- ▶ The first leap year of the Julian Calendar (as planned) 41BCE
- ▶ The first day of the Julian Calendar is 45BCE-01-01
- ▶ The last day of the Julian Calendar is 1582-10-4 (Thursday)
- ▶ The first day of the Gregorian Calendar is 1582-10-15 (Friday)
- ▶ In 1582, dates 10-5 through 10-14 didn't occur in Roman Catholic countries (it is named after Pope Gregory XIII)
- ▶ This change returned the equinoxes to the same days as the first year of the Julian Calendar
- ▶ Gregorian Calendar Leap Year Rule
Every year that is divisible by four is a leap year
Except for years that are divisible by 100
Unless they are divisible by 400, e.g., Y2K

The Julian Period, Proleptic Calendars and ISO 8601

- ▶ Proleptic Calendars adopt the rules of a Calendar and they run backwards in time before that Calendar starts
- ▶ The Julian Period starts with Proleptic Julian Date 4713BCE-01-01 and covers all known recorded human history
- ▶ ISO 8601 dictates a Proleptic Gregorian Calendar
- ▶ The Julian and Gregorian Calendars have no year zero
- ▶ ..., 2BCE, 1BCE, 1CE, 2CE, ...
- ▶ But ISO 8601 dictates the following definition of year

Year	ISO 8601	Julian
	-4712	4713BCE
	⋮	⋮
	-1	2BCE
	0	1BCE
	1	1CE
	⋮	⋮

A brief overview of R

- ▶ Interpreted language (very slow as opposed to compiled)
- ▶ Interpreter written in C and Fortran
- ▶ Considered a multi-paradigm language
- ▶ For example, it is considered to be object-oriented and functional (as well as other paradigms)
- ▶ Object-oriented and vectorized for user-friendliness: vectors and matrices are natural choices for *objects*
- ▶ “Everything is an object”
- ▶ Dynamically Typed (as opposed to Statically Typed)
implicitly typed by first usage
rather than explicitly by definition/declaration
- ▶ Provides multi-threading by *forking* **EXCEPT on Windows**
Windows does not have this capability so no multi-threading
An odd choice since MS was an early licensee of UNIX
Pioneering the port of UNIX to many CPUs as XENIX
And MS-DOS was a “dumbed-down” Unix-like system
If GNU Linux can clone forking, then why not MS?

A brief overview of R

- ▶ Much R functionality is written in R itself with calls to compiled C/C++/Fortran as needed (very fast)
- ▶ Mainly, these are R *add-on* packages as explained in the R FAQ; see section 5.1 in <https://cran.r-project.org/doc/FAQ/R-FAQ.html>
- ▶ There are two groups of packages distributed with R
- ▶ The 14 so-called *base* packages: `base`, `compiler`, `datasets`, `grDevices`, `graphics`, `grid`, `methods`, `parallel`, `splines`, `stats`, `stats4`, `tcltk`, `tools` and `utils`

```
> installed.packages(priority="base")
```
- ▶ The 15 so-called *recommended* packages: `KernSmooth`, `MASS`, `Matrix`, `boot`, `class`, `cluster`, `codetools`, `foreign`, `lattice`, `mgcv`, `nlme`, `nnet`, `rpart`, `spatial` and `survival`

```
> installed.packages(priority="recommended")
```
- ▶ Comprehensive R Archive Network: 18558 packages <https://cran.r-project.org/web/packages/index.html>
Task Views <https://cran.r-project.org/web/views>

Introduction to R by Venables and Smith: R-intro

- ▶ Without emacs, to launch R: `$ module load R/4.0.4; R`
- ▶ With emacs, launch the latest version of R: `M-x R-4.0.4` or simply `M-x R`
- ▶ The R quit command: `q()`
- ▶ Getting help: `> ?q` or `> help("q")`
- ▶ In the `*R:~*` at the `>`, type `??` for **help on help**
- ▶ Almost everything is an object (except keywords)
- ▶ To see the definition of `q`, type `> q`
- ▶ To search R help and installed packages help
type `> help.start()`
VERY SLOW IF LOTS OF PACKAGES ARE INSTALLED
- ▶ Double quotes vs. single quotes: the docs shows double quotes in most places, but single quotes seem to work fine and they are easier to type: I will try to use double in my notes

Section 1: Storing and re-loading objects

- ▶ Names contain alphanumeric characters, underscore or period
- ▶ If they start with period, the second character must be a letter
- ▶ A case-sensitive language, i.e., distinct variables `a` and `A`
- ▶ The `source` function is convenient for loading objects like data, functions, etc.
- ▶ `source("lecture2.R")`
- ▶ However, re-creating some objects (like large data sets or complex analyses) might be very time-consuming
- ▶ DON'T DO THIS: R-intro describes saving R objects in `.RData` which is automatically re-loaded when you re-launch R in that directory: IT IS TOO EASY TO FORGET/MISUSE
- ▶ Save the objects that you want to keep and re-load as needed
- ▶ `saveRDS(object, "object.rds")`
- ▶ `object = readRDS("object.rds")`
- ▶ R loads all objects into memory so be cognizant of their size
- ▶ `object.size(object)`
- ▶ `rm(object); gc()`
- ▶ `?gc`

Section 1: View-/edit-ing objects

- ▶ You can view an object just by typing its name at the prompt:
`> object`
- ▶ However, for large objects, this is not very user-friendly
- ▶ For large objects, it might be better to view/edit them
- ▶ If you just want to view (and NOT edit):
`> sink("object.Rout"); object; sink()`
`C-x C-f object.Rout`
- ▶ To edit, start an emacsclient by `M-x server-start`
- ▶ `> sink("/dev/null") ## shut off the R console`
- ▶ `> ls() ## notice that it produces NO output`
- ▶ `> edit(object, "object.R")`
- ▶ When you are done: `C-x #`
- ▶ `> sink() ## restore the R console`
- ▶ `> ls() ## now produces output as usual`
- ▶ N.B. the file `object.R` remains for your viewing

Section 2: R is naturally vectorized

- ▶ All atomic types are vectors of length zero or more: numeric, logical, character, complex and raw
- ▶ Roughly in that order of importance
- ▶ `length` function returns their current length
- ▶ Vectors can easily expand by allocating beyond it
- ▶ `> x = 1` creates a vector of length 1
- ▶ `> (x = 1)` parentheses echo assignment values
- ▶ Like `> x = 1; x`
- ▶ Or to echo any expression `> 1/x`
- ▶ `> x[2]=0` expands to length of 2
- ▶ `> x = c(10.4, 5.6, 3.1, 6.4, 21.7)` creates a vector
- ▶ `> ?c`
- ▶ see `lecture3.R`

Section 2.3: Sequences

- ▶ Integer sequences: `1:30` same as `c(1, 2, ..., 29, 30)`
- ▶ General sequences: `seq(-5, 5, by=0.2)` same as `c(-5.0, -4.8, -4.6, ..., 4.6, 4.8, 5.0)`
- ▶ See `> ?seq` for all of the possibilities

Section 2.4: Logical vectors

- ▶ Three possible values: **TRUE**, **FALSE** and NA (BEWARE!)
- ▶ The comparison operators: `<`, `<=`, `>`, `>=`, `==`, `!=`
- ▶ Two equal signs for equality like C
- ▶ Combining expressions with `|`, `||`, `&`, `&&`
- ▶ N.B. the difference between ORs `|` and `||`
and the difference between ANDs `&` and `&&`
- ▶ `||` and `&&` SHOULD ONLY BE USED WITH SCALARS
they stop evaluating once the result is known
so very useful for conditional expressions
`FALSE || TRUE || FALSE`
`TRUE && FALSE && TRUE`
- ▶ `|` and `&` are for vector operations (similar to addition/etc.)
- ▶ Functions that return a scalar from vectors: `any` and `all`
- ▶ `TRUE` can be coerced to 1 and `FALSE` can be coerced to 0
- ▶ see `lecture3.R`

Section 2.6: Character vectors

- ▶ Working with character strings is NOT one of R's strengths
- ▶ But this is true of most languages that statistician's use
- ▶ The focus of statistics is largely on numerical computing
- ▶ Many functions inherited from the AWK language (mainly New AWK or NAWK) from the UNIX family (cloned by the GNU as GAWK)
- ▶ For example, see `sub`, `gsub` and `substr`
- ▶ Three functions commonly used (NOT inherited from AWK): `paste0`, `paste`, `nchar` and `sprintf`
- ▶ see `lecture3.R`

Section 2.7: Vector indices

- ▶ Integer indices: `x[1:3]` returns the first three items
- ▶ If there are fewer, then returns NA
- ▶ `x[0]` returns nothing
- ▶ `x[-1]` returns everything EXCEPT `x[1]`
- ▶ Logical indices: `x[!is.na(x)]`
- ▶ BEWARE: due to coercion of complex expressions you can accidentally request TRUE as index 1 and FALSE as index 0 (which is nothing): can be very difficult to debug because it happens without warning
- ▶ You can protect yourself by enclosing all logical index expressions within `which` function. See `?which`
- ▶ `x[which(!is.na(x))]`

Section 3: Object attributes

- ▶ All objects have a *mode*: `> mode(x)`
- ▶ But mainly of interest for atomic vectors
- ▶ Some objects have a *class*: `> class(x)`
- ▶ A capable function for object structure: `> str(x)`
- ▶ `as.` functions for manual coercion between types
- ▶ `> as.integer(x)`
- ▶ `> as.character(x)`
- ▶ All of an object's attributes, if any, can be listed:
`attributes(x)`
- ▶ They can be extracted: `attr(x, "names")`
- ▶ And set: `attr(x, "names")=letters[1:5]`

Appendices

- ▶ Appendix A lists a lot of interesting functions that we will learn about
- ▶ But, `attach` is generally frowned upon today
- ▶ A short version of the help provided in Appendix B.1 is generated by `$ R --help`

Appendix B.4: Scripting with R

- ▶ It is not about “Scripting” per se
- ▶ Rather, asynchronous/background/batch R jobs
i.e., NOT interactive
- ▶ It does provide useful information (which is why it is assigned)
- ▶ Particularly, the function `commandArgs` and `stdin` discussion
- ▶ Generally, I do something like the following
- ▶ the standard Unix way of doing things with re-direction
building on what is shown in Appendix B.1
- ▶ `stdin` is coming from `trees.R`
- ▶ `stdout` and `stderr` are going to `trees.Rout`

```
nohup R --no-save < trees.R >& trees.Rout &
```